

# Python and Data Science (Live Coding)

A. Richards

09.13.2017



# Why Data Science?

Well because of the **data** and because of the **science**

- Data science means different things to different people
  - Data engineering → data vis → predictive modeling
- Statistics, machine-learning, databases, web-development
- We are in an era of unprecedented data growth
  - And so few know how to effectively use data to gain insight
- Science is the proponent of truth through the communication of evidence

# Why Data Science?

Well because of the **data** and because of the **science**

- Data science means different things to different people
  - Data engineering → data vis → predictive modeling
- Statistics, machine-learning, databases, web-development
- We are in an era of unprecedented data growth
  - And so few know how to effectively use data to gain insight
- Science is the proponent of truth through the communication of evidence

There are jobs as well!

# The essential data science toolkit

## Subject area mastery

- SQL and noSQL databases
- Associative statistics and hypothesis testing
- Unsupervised and supervised learning
- Data visualization
- Data products


## Programming mastery

Expert level proficiency in language that is useful for data science

# Why Python?

- There are many languages of data science (Python, R,...)
- Ecosystem (NumPy, matplotlib, pandas)
- Readability, Flexibility
- Glue language
- Object-oriented and functional


# What is NLP?



- Conversational Agents
  - Siri, Cortana, Google Now, Alexa
  - Talking to your car
  - Communicating with robots
- Machine Translation
  - Google Translate
  - Google's Neural Machine Translation
- Speech Recognition, Speech Synthesis
- Lexical Semantics, Sentiment Analysis
- Dialogue Systems, Question Answering

# NLP and AI

The ultimate goal of NLP is to fill the gap how the humans communicate (natural language) and what the computer understands (machine language).




## Why Deep Learning Needed in NLP

- It uses a rule-based approach that represents Words as 'One-Hot' encoded vectors.
- Traditional method focuses on syntactic representation instead of semantic representation.
- Bag of words - classification model is unable to distinguish certain contexts.


<http://www.datasciencecentral.com/profiles/blogs/overview-of-artificial-intelligence-and-role-of-natural-language>



# Classical NLP



# Deep Learning



What is the meaning of this sentence?

I made her duck.

# Challenges

## Ambiguity

What does it mean when we say: 'I made her duck'

- I cooked waterfowl for her
- I cooked waterfowl belonging to her
- I created the (papier mache?) duck she owns
- I caused her to quickly lower her head or body
- I waved my magic wand and turned her into undifferentiated waterfowl

## Other examples

- 'Court to try shooting defendant'
- 'Hospitals are sued by seven foot doctors'


This problem of determining which sense was meant by a specific word is formally known as **word sense disambiguation**

Other challenges include:

- Part of speech tagging
- Syntactic disambiguation (The I made her duck example)

# Demo 1

## Tutorial Example



## Demo 2

Advanced Example

# Overview

- 1 Version control based workflow
- 2 Jupyter notebook
- 3 Data types (lists, arrays, data frames)
- 4 EDA - working with data frames, plotting
- 5 Working with text (latent topics, clustering)
- 6 Data visualization

Week	Topic
0	Python workshop
1	Programming and SQL
2	Probability and Statistics
3	Linear Models
4	Supervised Learning
5	NLP, Business analytics
6	Unsupervised Learning
7	Break Week
8	Big Data / Data Engineering
9	Special Topics / Case Studies
9	Project
10	Project
11	Project
12	Career services and Special topics

There are also part-time courses and free meet-ups...